

Sentimental Analysis on Tweets

Iro Moumoulidou

Technical University Of Crete

June 23, 2016

Introduction

- Main Object: Classify tweets as positive/negative
- Build a Classifier with Supervised Learning

Our Data Set

- Data set retrieved from an in-class competition organized at Kaggle
- Set of tweets already classified as negative or positive
- 84.400 of them were used to train our classifier, 8.420 to test it

Tweet Representation

- Every Tweet is a **Bag of Words** (positive or negative respectively)
- Union of same sentiment *bags of words*, form the terms of each class
- However, do we need every single one of these terms?

Detection Of Uninformative Terms

- Characteristics of tweets: @username, #feelings , urls
- Stopwords
- What about terms that appear in both classes?

Detection Of Uninformative Terms

- For a term x of both classes a metric is defined as:

$$ratio = \frac{\min(P_x, N_x)}{\max(P_x, N_x)}$$

- If this ratio is close to one, occurrences of the term are almost equal
- A good threshold should be found for discarding these terms

Detection Of Uninformative Terms

- Another idea: Select the most representative terms-words for each class
- Metric used : tf as in Information Retrieval
- Term-frequency(tf) in a class is defined as:

$$tf_c = \frac{\text{frequency of term in class } c}{\text{maximum frequency of class } c}$$

- Again a good threshold should be found for discarding irrelative terms

Multinomial Naive Bayes

- Model used to train our classifier is Multinomial Naive Bayes
- Given a tweet t , the classifier should choose the class that maximizes the probability $p(c|t)$
- Formally we get:

$$\begin{aligned}c_{NB} &= \operatorname{argmax}_c p(c|t) \\&= \operatorname{argmax}_c \frac{p(t|c)p(c)}{p(t)} \\&= \operatorname{argmax}_c p(t|c)p(c)\end{aligned}$$

Multinomial Naive Bayes

- However a tweet is a set of independent terms, so we get:

$$\begin{aligned}c_{NB} &= \underset{c}{\operatorname{argmax}} p(t|c)p(c) \\&= \underset{c}{\operatorname{argmax}} p(t_1, t_2, \dots, t_n|c)p(c) \\&= \underset{c}{\operatorname{argmax}} p(t_1|c)p(t_2|c) \dots p(t_n|c)p(c)\end{aligned}$$

- The probability of each term t_i is defined as:

$$p(t_i|c) = \frac{\text{frequency of term } t_i \text{ in class } c + 1}{\text{sum of all frequencies in class } c + |D|}$$

- $|D|$: number of distinct words in both classes

Multinomial Naive Bayes

- Finally, the a-priori probability of each class is defined as:

$$p(c) = \frac{\text{num of processed tweets in class } c}{\text{total num of processed tweets}}$$

Preprocessing Of Test Tweets

Every test-tweet has undergone the same preprocessing(stopword removal etc) as the training tweets.

Experiments

- Quick overview of our two thresholds:

$$ratio = \frac{\min(P_x, N_x)}{\max(P_x, N_x)}$$

$$tf_c = \frac{\text{frequency of term in class } c}{\text{maximum frequency of class } c}$$

- First, we set a threshold for ratio metric and discard every term above it.
- Then, among the remaining tweets in each class those with tf less than the second threshold are discarded

Experiments

- The error results of our classifier for various thresholds are:

$t1 \setminus tf$	0.1	0.2	0.5	0.7
0.9	0.32	0.36	0.46	0.464
0.7	0.33	0.37	0.43	0.464
0.5	0.34	0.38	0.45	0.458

Table 1: Errors for Various Thresholds

- As we decrease the $t1$, we increase the number of terms chosen for deletion according to the first filter.
- We observe that the error in this case does not change much.

Experiments

- As we increase the tf threshold, the error deteriorates as we choose to keep less and less terms.
- So the tf thresholds seems to determine the final error, and the more terms are included in the feature vector of each class, the better performance we get.



Thank You!

A handwritten message in green ink on a light gray background. The text "Thank You!" is written in a cursive, friendly style. Below the word "Thank" is a simple smiley face drawn with two dots for eyes and a curved line for a mouth.